

## EXPERIMENT NO. 1

**TITLE :** Perform EDA and Visualise the relationships for a given dataset.

**CODE :**

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the Dataset
data = pd.read_csv("Housing.csv")
print(data.info())
print(data.describe())

# Check for Missing Values
print("Missing Values:\n", data.isnull().sum())

# Univariate Analysis: Distribution of Prices
sns.histplot(data['price'], kde=True)
plt.title("Distribution of House Prices")
plt.xlabel("Price")
plt.ylabel("Frequency")
plt.show()

# Distribution of Area
sns.histplot(data['area'], kde=True, color='orange')
plt.title("Distribution of House Area")
plt.xlabel("Area")
plt.ylabel("Frequency")
plt.show()

# Bar Plot: Furnishing Status
sns.countplot(x="furnishingstatus", data=data, palette="Set2")
```

```
plt.title("Furnishing Status Count")
plt.xlabel("Furnishing Status")
plt.ylabel("Count")
plt.show()
```

```
# Multivariate Analysis: Price vs Area
sns.scatterplot(x="area", y="price", hue="furnishingstatus", data=data, palette="cool")
plt.title("Price vs Area")
plt.xlabel("Area")
plt.ylabel("Price")
plt.show()
```

```
# Correlation Heatmap
sns.heatmap(data.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```

```
# Price vs Number of Bedrooms
sns.boxplot(x="bedrooms", y="price", data=data, palette="Set3")
plt.title("Price vs Number of Bedrooms")
plt.xlabel("Number of Bedrooms")
plt.ylabel("Price")
plt.show()
```

## OUTPUT :


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad            545 non-null   object
6   guestroom           545 non-null   object
7   basement            545 non-null   object
8   hotwaterheating     545 non-null   object
9   airconditioning     545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
None
```

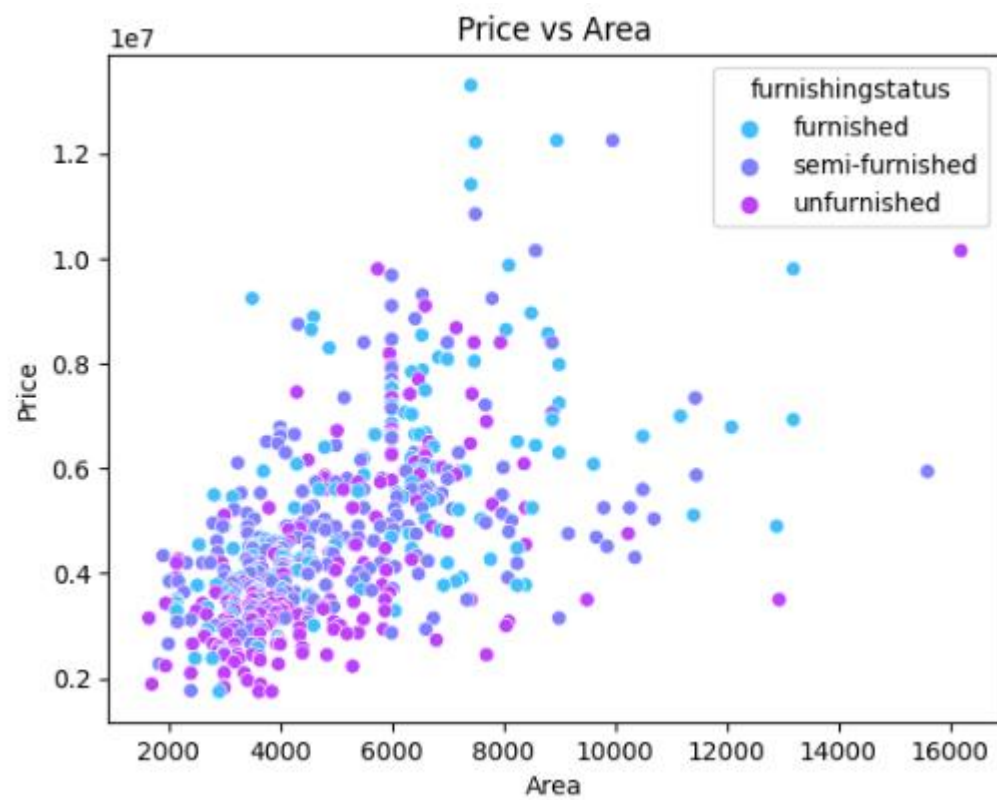
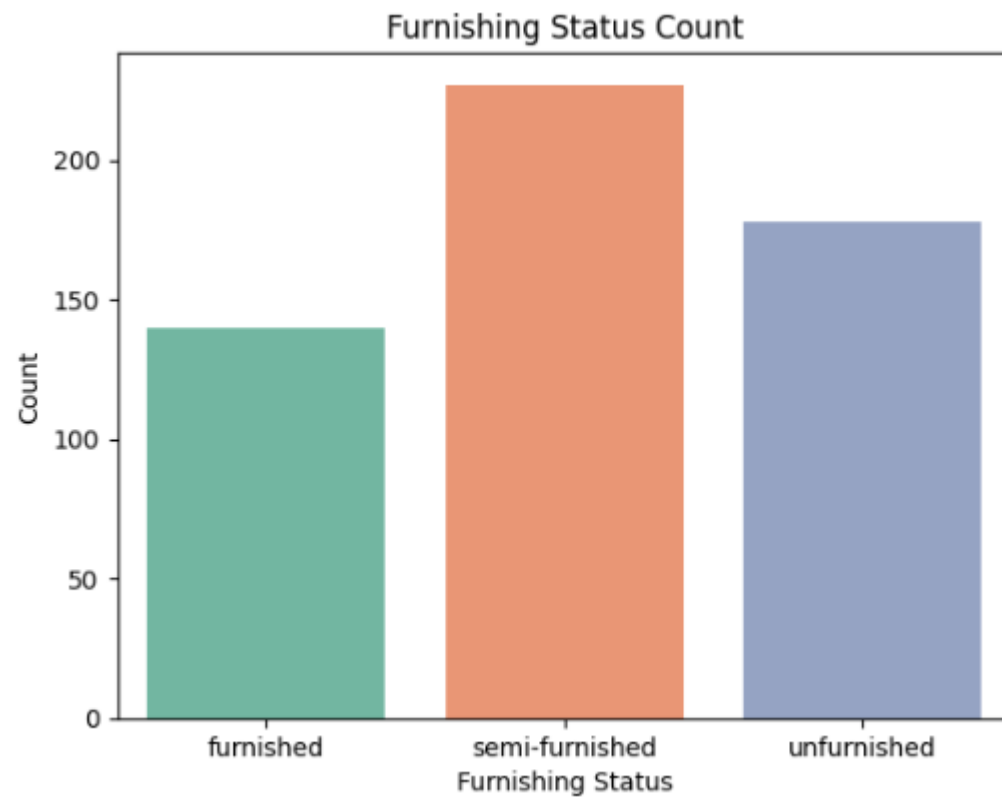
	price	area	bedrooms	bathrooms	stories
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

```
Missing Values:
price      0
area      0
bedrooms   0
bathrooms  0
stories    0
mainroad   0
guestroom  0
basement   0
hotwaterheating  0
airconditioning  0
parking    0
prefarea   0
furnishingstatus  0
dtype: int64
```





## EXPERIMENT NO. 2

**TITLE :** Perform EDA and Visualise the relationships and also train a Linear regression model and evaluate model's performance for Housing Price dataset.

**CODE :**

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load Dataset
data = pd.read_csv("User_Data.csv")
print(data.info())

# EDA: Visualize the relationships
sns.scatterplot(x="EstimatedSalary", y="Purchased", hue="Gender", data=data)
plt.title("Relationship Between Salary and Purchase Decision")
plt.show()

# Preprocessing
data = pd.get_dummies(data, columns=["Gender"], drop_first=True)
X = data[["Age", "EstimatedSalary", "Gender_Male"]]
y = data["Purchased"]

# Split Data into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Linear Regression Model
model = LinearRegression()
model.fit(X_train, y_train)

# Make Predictions
```

```

y_pred = model.predict(X_test)

# Evaluate Model

rmse = mean_squared_error(y_test, y_pred, squared=False)

r2 = r2_score(y_test, y_pred)

print("RMSE:", rmse)

print("R² Score:", r2)

```

## OUTPUT :

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased             400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
None

```



```

RMSE: 0.315190846926146
R² Score: 0.5633174945669398

```

### EXPERIMENT NO. 3

**TITLE :** Perform EDA and Visualise the relationships and also train a Ridge regression model and evaluate model's performance for given dataset.

**CODE :**

```
# Import necessary libraries

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset

df = pd.read_csv("Housing.csv")
print(df.shape)

# EDA

sns.pairplot(df)

plt.show()

sns.heatmap(df.corr(), annot=True, cmap="coolwarm")

plt.show()

# Preprocessing

df = pd.get_dummies(df, drop_first=True)

X = df.drop("price", axis=1)

y = df["price"]

# Standardization
```



```

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

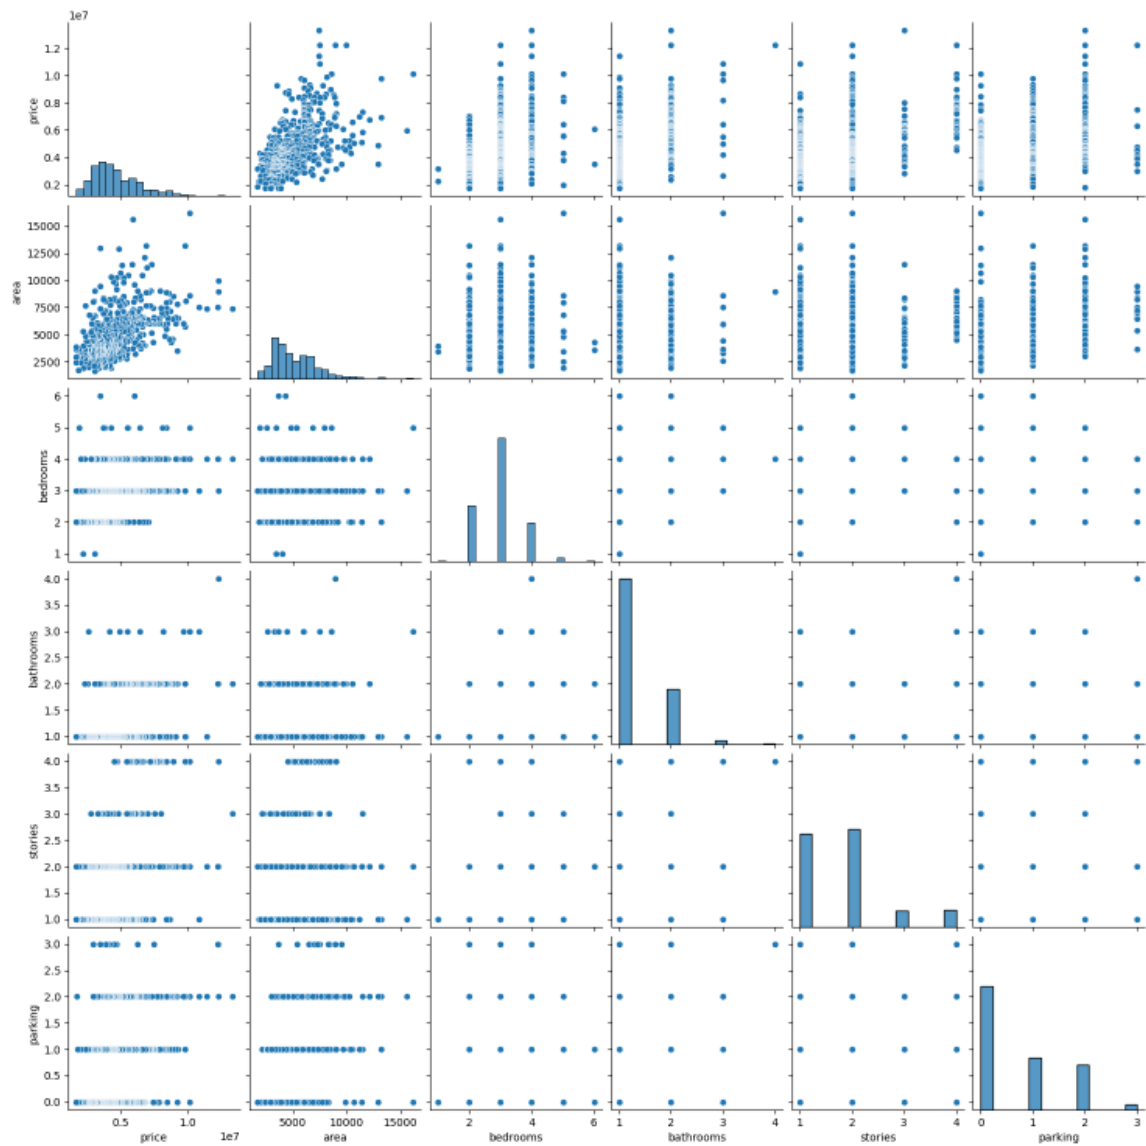

# Model Training and Evaluation
ridge = Ridge(alpha=1.0) # Try different alpha ( $\lambda$ ) values
ridge.fit(X_train, y_train)
y_pred = ridge.predict(X_test)



# Performance Metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")


# Residuals Plot
residuals = y_test - y_pred
sns.histplot(residuals, kde=True)
plt.title("Residuals Distribution")
plt.show()


```

OUTPUT :





Mean Squared Error: 1754768938809.1035  
R-squared: 0.6528351861223265



#### EXPERIMENT NO. 4

**TITLE :** Perform EDA and Visualise the relationships and also train a Lasso regression model and evaluate model's performance for given dataset.

**CODE :**

```
# Import necessary libraries

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset

df = pd.read_csv("Housing.csv")
print(f"Dataset Shape: {df.shape}")

# EDA

sns.pairplot(df)

plt.show()

sns.heatmap(df.corr(), annot=True, cmap="coolwarm")

plt.show()

# Preprocessing

df = pd.get_dummies(df, drop_first=True)

X = df.drop("price", axis=1)

y = df["price"]

# Standardization

from sklearn.preprocessing import StandardScaler
```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

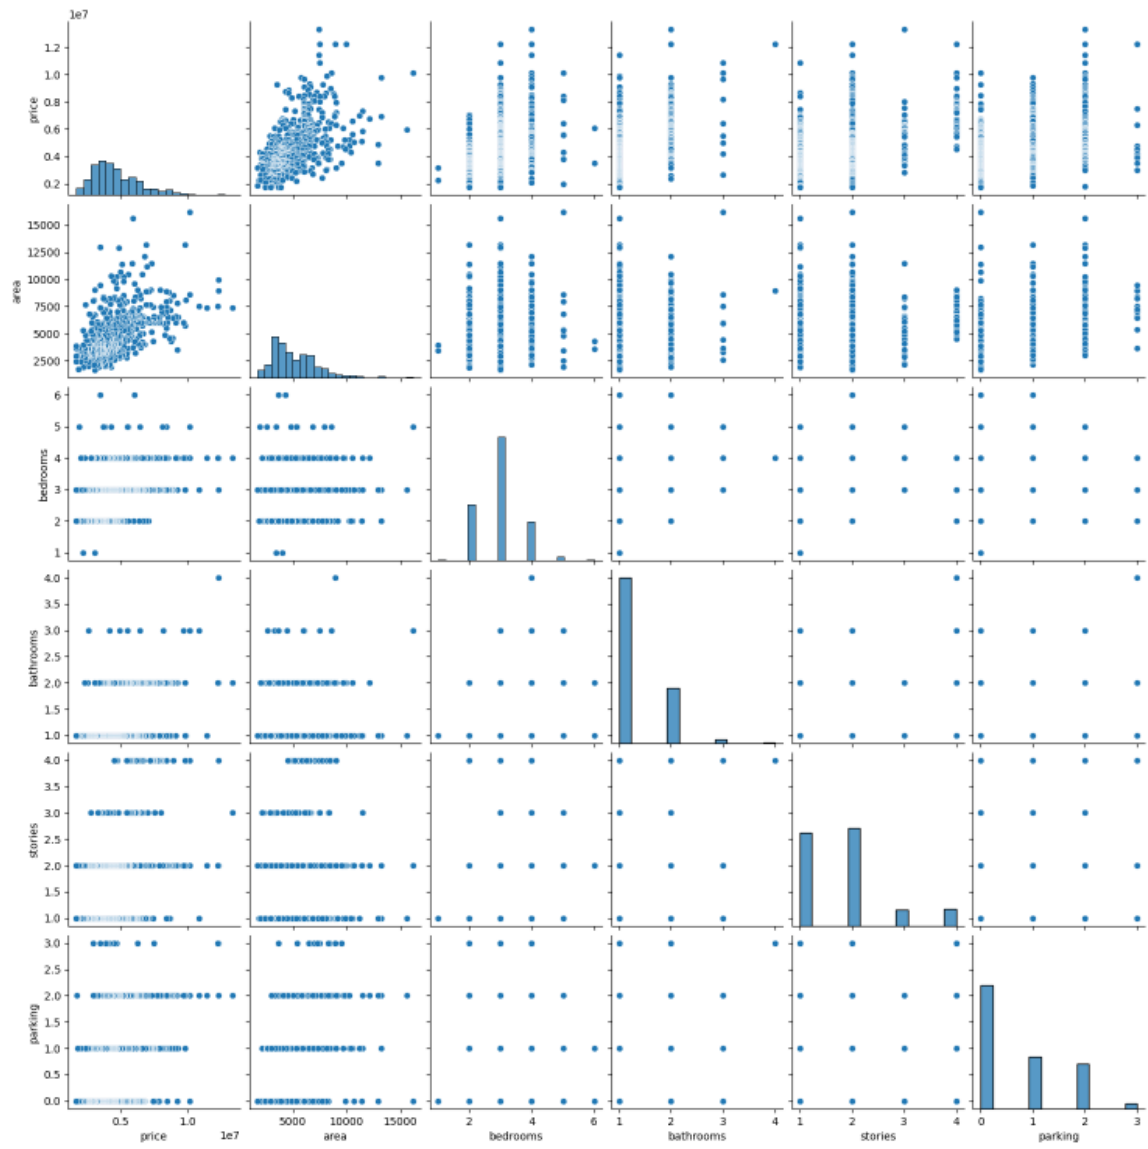
# Model Training and Evaluation
lasso = Lasso(alpha=1.0) # Experiment with different alpha ( $\lambda$ ) values
lasso.fit(X_train, y_train)
y_pred = lasso.predict(X_test)


# Performance Metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")

# Residual Plot
residuals = y_test - y_pred
sns.histplot(residuals, kde=True)
plt.title("Residuals Distribution")
plt.show()


```

**OUTPUT :**





Mean Squared Error: 1754768938809.1035  
R-squared: 0.6528351861223265





## EXPERIMENT NO. 5

**TITLE :** Perform EDA and Visualise the relationships and also train a Logistic regression model and evaluate model's performance for Iris dataset.

**CODE :**

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load Dataset
from sklearn.datasets import load_iris
data = load_iris(as_frame=True)
df = data.frame

# EDA: Pair plot
sns.pairplot(df, hue='target', vars=['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)'])
plt.show()

# Split Data
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Logistic Regression
model = LogisticRegression()
model.fit(X_train, y_train)

# Predictions and Evaluation
```

```

y_pred = model.predict(X_test)

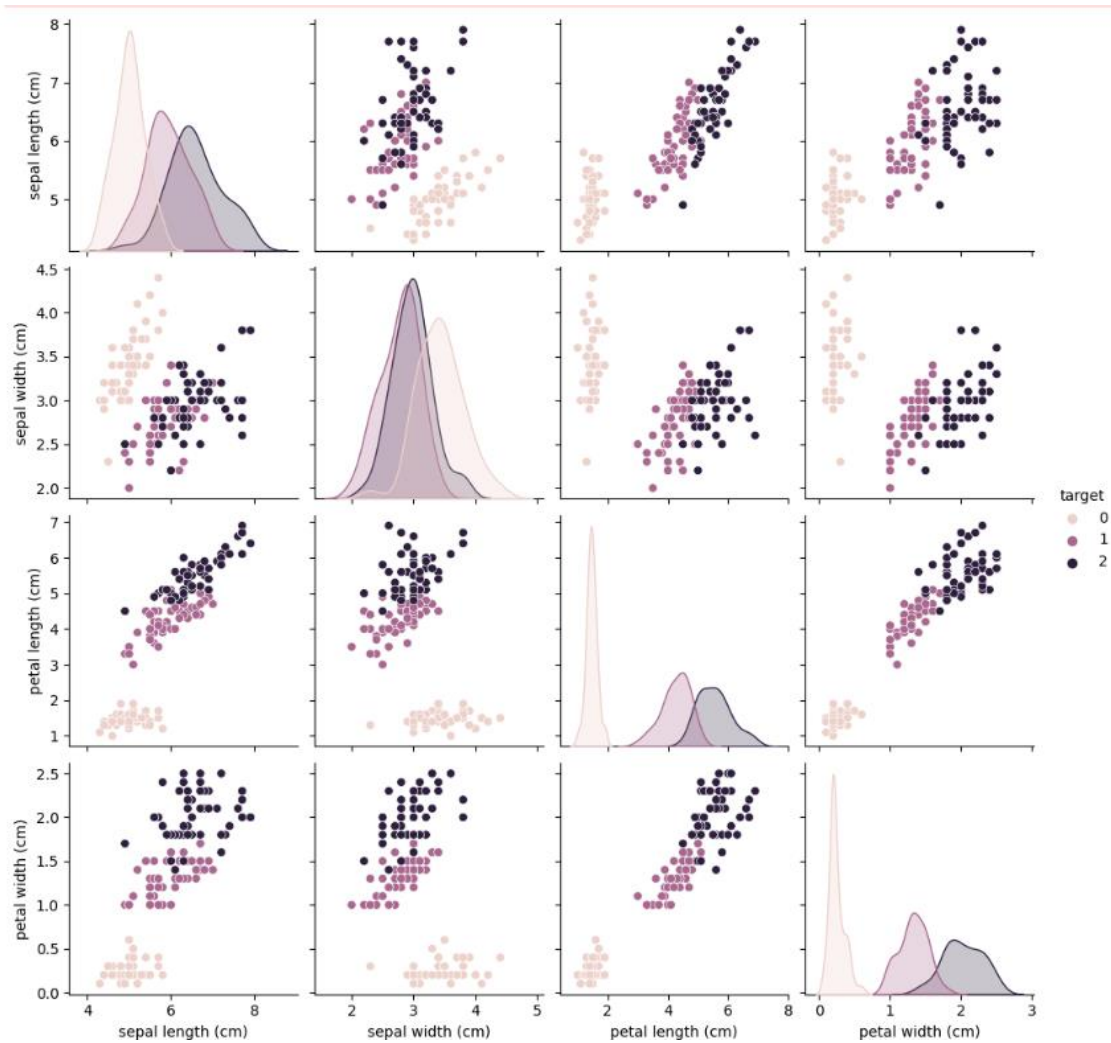
print("Accuracy:", accuracy_score(y_test, y_pred))

print("Classification Report:\n", classification_report(y_test, y_pred))

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

## OUTPUT :



```

Accuracy: 1.0
Classification Report:
              precision    recall  f1-score   support

     0           1.00       1.00       1.00        10
     1           1.00       1.00       1.00         9
     2           1.00       1.00       1.00        11

 accuracy          1.00          1.00          1.00        30
  macro avg          1.00          1.00          1.00        30
weighted avg          1.00          1.00          1.00        30

Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]

```

## EXPERIMENT NO. 6

**TITLE :** Perform EDA and Visualise the relationships and also train a Logistic regression model and evaluate model's performance for User dataset.

**CODE :**

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Load Dataset
data = pd.read_csv('User_Data.csv')

# EDA: Check dataset info
print(data.info())

# Scatter Plot
sns.scatterplot(x='Age', y='EstimatedSalary', hue='Purchased', data=data)
plt.title('Age vs Estimated Salary Colored by Purchase')
plt.show()

# Data Preprocessing
data['Gender'] = data['Gender'].map({'Male': 0, 'Female': 1}) # Encode Gender
X = data[['Gender', 'Age', 'EstimatedSalary']]
y = data['Purchased']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Logistic Regression
model = LogisticRegression()
```

```
model.fit(X_train, y_train)
```

```
# Predictions and Evaluation
```

```
y_pred = model.predict(X_test)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

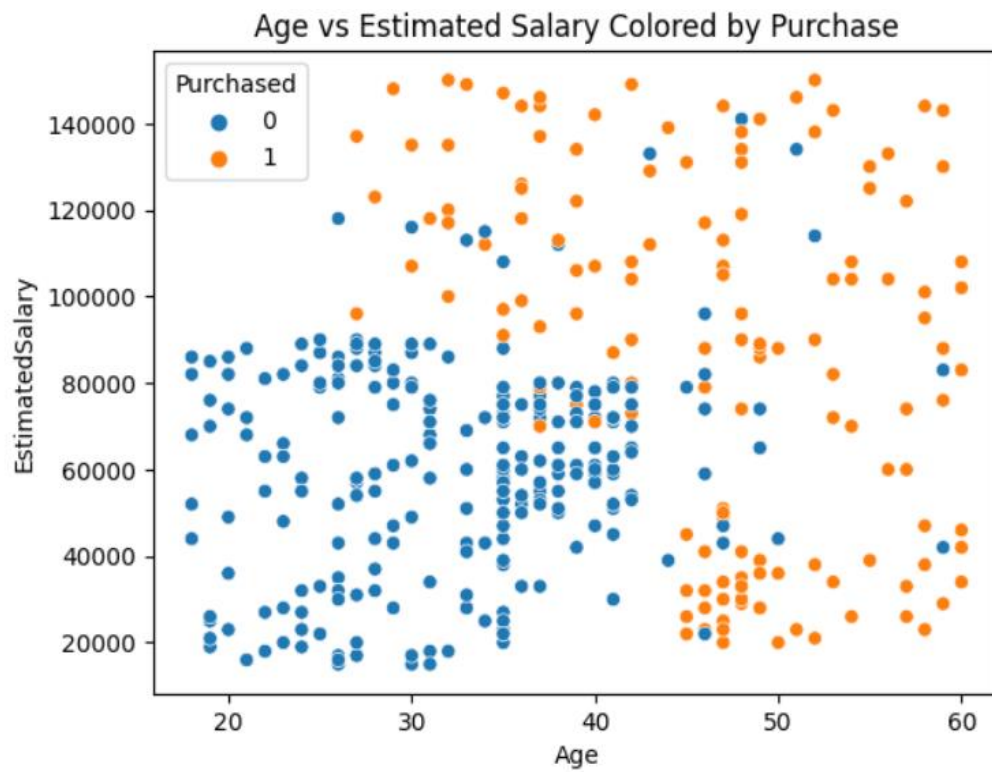
```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

**OUTPUT :**

---

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID                400 non-null   int64
1   Gender                 400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased              400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
None
```



Accuracy: 0.65

Classification Report:

	precision	recall	f1-score	support
0	0.65	1.00	0.79	52
1	0.00	0.00	0.00	28
accuracy			0.65	80
macro avg	0.33	0.50	0.39	80
weighted avg	0.42	0.65	0.51	80

Confusion Matrix:

```
[[52  0]
 [28  0]]
```

## EXPERIMENT NO. 7

**TITLE :** Perform EDA and Visualise the relationships and also train a Naive Baye's Classifier model and evaluate model's performance for Wine dataset.

**CODE :**

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.naive_bayes import GaussianNB

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt


# Load Dataset

from sklearn.datasets import load_wine

data = load_wine(as_frame=True)

df = data.frame


# EDA: Visualizing Correlation Heatmap

sns.heatmap(df.corr(), annot=True, cmap='coolwarm')

plt.title('Correlation Heatmap')

plt.show()


# Splitting Features and Target

X = df.drop('target', axis=1)

y = df['target']


# Train-Test Split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Train Naive Bayes Classifier

model = GaussianNB()
```

```
model.fit(X_train, y_train)
```

```
# Predictions and Evaluation
```


```
y_pred = model.predict(X_test)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

**OUTPUT :**





Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	14
1	1.00	1.00	1.00	14
2	1.00	1.00	1.00	8
accuracy			1.00	36
macro avg	1.00	1.00	1.00	36
weighted avg	1.00	1.00	1.00	36

Confusion Matrix:

```
[[14  0  0]
 [ 0 14  0]
 [ 0  0  8]]
```

## EXPERIMENT NO. 8

**TITLE :** Perform EDA and Visualise the relationships and also train a Decision Tree Classifier model and evaluate model's performance for Titanic dataset.

**CODE :**

```
import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix


# Load Dataset

data = pd.read_csv('titanic.csv')


# EDA: Summary and Visualization

print(data.info())

sns.barplot(x='Sex', y='Survived', data=data)

plt.title('Survival Rate by Gender')

plt.show()


# Preprocessing: Handle missing values and encode categorical variables

data['Age'].fillna(data['Age'].median(), inplace=True)

data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)

data = pd.get_dummies(data, columns=['Sex', 'Embarked'], drop_first=True)


# Feature Selection

features = ['Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Sex_male', 'Embarked_Q', 'Embarked_S']

X = data[features]

y = data['Survived']


# Split Data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train Decision Tree Classifier
```

```
model = DecisionTreeClassifier(max_depth=5, random_state=42)
```

```
model.fit(X_train, y_train)
```

```
# Predictions and Evaluation
```

```
y_pred = model.predict(X_test)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```


```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

**OUTPUT :**

---

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null   int64
 1   Survived        891 non-null   int64
 2   Pclass         891 non-null   int64
 3   Name           891 non-null   object
 4   Sex            891 non-null   object
 5   Age           714 non-null   float64
 6   SibSp         891 non-null   int64
 7   Parch         891 non-null   int64
 8   Ticket        891 non-null   object
 9   Fare          891 non-null   float64
10   Cabin         204 non-null   object
11   Embarked      889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```



Accuracy: 0.7988826815642458

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.90	0.84	105
1	0.83	0.65	0.73	74
accuracy			0.80	179
macro avg	0.81	0.78	0.78	179
weighted avg	0.80	0.80	0.79	179

Confusion Matrix:

```
[[95 10]
 [26 48]]
```

## EXPERIMENT NO. 9

**TITLE :** Perform EDA and Visualise the relationships and also train a KNN model and evaluate model's performance for Digit dataset.

**CODE :**

```
import numpy as np

import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.datasets import load_digits

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report


# Load Digits Dataset

digits = load_digits()

X = digits.data

y = digits.target

print(f"Dataset Shape: {X.shape}")


# Visualize a Sample Digit

plt.gray()

plt.matshow(digits.images[0])

plt.title(f"Digit: {digits.target[0]}")

plt.show()


# Normalize the Data

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# Split into Training and Testing Sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)


# Train KNN Model
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Make Predictions
y_pred = knn.predict(X_test)

# Evaluate Model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

## OUTPUT :

Dataset Shape: (1797, 64)  
<Figure size 640x480 with 0 Axes>



Accuracy: 0.975

Confusion Matrix:

```
[[33  0  0  0  0  0  0  0  0  0]
 [ 0 28  0  0  0  0  0  0  0  0]
 [ 0  0 33  0  0  0  0  0  0  0]
 [ 0  0  1 33  0  0  0  0  0  0]
 [ 0  0  0  0 46  0  0  0  0  0]
 [ 0  0  0  0  0 45  1  0  0  1]
 [ 0  0  0  0  0  0 35  0  0  0]
 [ 0  0  0  0  0  1  0 32  0  1]
 [ 0  0  0  0  0  0  0  0 30  0]
 [ 0  0  0  1  1  1  0  0  1 36]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	33
1	1.00	1.00	1.00	28
2	0.97	1.00	0.99	33
3	0.97	0.97	0.97	34
4	0.98	1.00	0.99	46
5	0.96	0.96	0.96	47
6	0.97	1.00	0.99	35
7	1.00	0.94	0.97	34
8	0.97	1.00	0.98	30
9	0.95	0.90	0.92	40
accuracy			0.97	360
macro avg	0.98	0.98	0.98	360
weighted avg	0.98	0.97	0.97	360

## EXPERIMENT NO. 10

**TITLE :** . Perform EDA and Visualise the relationships and also train a K-Means regression model and evaluate model's performance for Income dataset.

**CODE :**

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

# Load Dataset
data = pd.read_csv("User_Data.csv")
print(data.info())

# EDA: Visualize EstimatedSalary and Age distributions
sns.scatterplot(x="Age", y="EstimatedSalary", hue="Purchased", data=data)
plt.title("Scatterplot: Age vs Estimated Salary Colored by Purchase")
plt.show()

# Preprocessing
data = pd.get_dummies(data, columns=["Gender"], drop_first=True)
scaler = StandardScaler()
data[["Age", "EstimatedSalary"]] = scaler.fit_transform(data[["Age", "EstimatedSalary"]])

# Apply K-Means Clustering
kmeans = KMeans(n_clusters=3, random_state=42)
data["Cluster"] = kmeans.fit_predict(data[["Age", "EstimatedSalary"]])

# Visualize Clusters
sns.scatterplot(x="Age", y="EstimatedSalary", hue="Cluster", data=data, palette="viridis")
plt.title("Clusters Visualized: Age vs Estimated Salary")
```



```
plt.show()
```

```
# Evaluate Model
```


```
print("Cluster Centroids:\n", kmeans.cluster_centers_)
```


```
print("Intra-Cluster Variance (WCSS):", kmeans.inertia_)
```

```
print("Silhouette Score:", silhouette_score(data[["Age", "EstimatedSalary"]], data["Cluster"]))
```

**OUTPUT :**

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 400 entries, 0 to 399  
Data columns (total 5 columns):  
#   Column          Non-Null Count  Dtype    
---  ---            
0   User ID         400 non-null   int64    
1   Gender          400 non-null   object   
2   Age             400 non-null   int64    
3   EstimatedSalary 400 non-null   int64    
4   Purchased       400 non-null   int64    
dtypes: int64(4), object(1)  
memory usage: 15.8+ KB  
None
```





Cluster Centroids:

```
[[ 0.7991185 -0.40988428]
```

```
[ 0.46159632 1.54420291]
```

```
[-0.83200206 -0.38911019]]
```

Intra-Cluster Variance (WCSS): 323.8667384794876

Silhouette Score: 0.36205845294545275

## EXPERIMENT NO. 11

**TITLE :** Perform EDA and Visualise the relationships and also train a K-Means regression model and evaluate model's performance for Iris dataset.

**CODE :**

```
import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.datasets import load_iris

from sklearn.preprocessing import StandardScaler

from sklearn.cluster import KMeans

from sklearn.metrics import silhouette_score

import numpy as np


# Load Dataset

iris = load_iris()

data = pd.DataFrame(iris.data, columns=iris.feature_names)

data['species'] = iris.target

print(data.head())


# EDA: Pair Plot

sns.pairplot(data, hue='species', vars=iris.feature_names)

plt.show()


# Data Preprocessing: Standardize Features

scaler = StandardScaler()

data_scaled = scaler.fit_transform(data[iris.feature_names])


# Determine Optimal k using Elbow Method

wcss = []

for k in range(1, 11):

    kmeans = KMeans(n_clusters=k, random_state=42)
```

```

kmeans.fit(data_scaled)
wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS')
plt.show()

# Train K-Means Model
kmeans = KMeans(n_clusters=3, random_state=42)
data['cluster'] = kmeans.fit_predict(data_scaled)

# Visualize Clusters
sns.scatterplot(x=data[iris.feature_names[0]], y=data[iris.feature_names[2]], hue=data['cluster'],
palette='viridis')
plt.title('K-Means Clusters')
plt.xlabel('Sepal Length')
plt.ylabel('Petal Length')
plt.show()

# Evaluate Clustering
silhouette_avg = silhouette_score(data_scaled, data['cluster'])
print("Silhouette Score:", silhouette_avg)

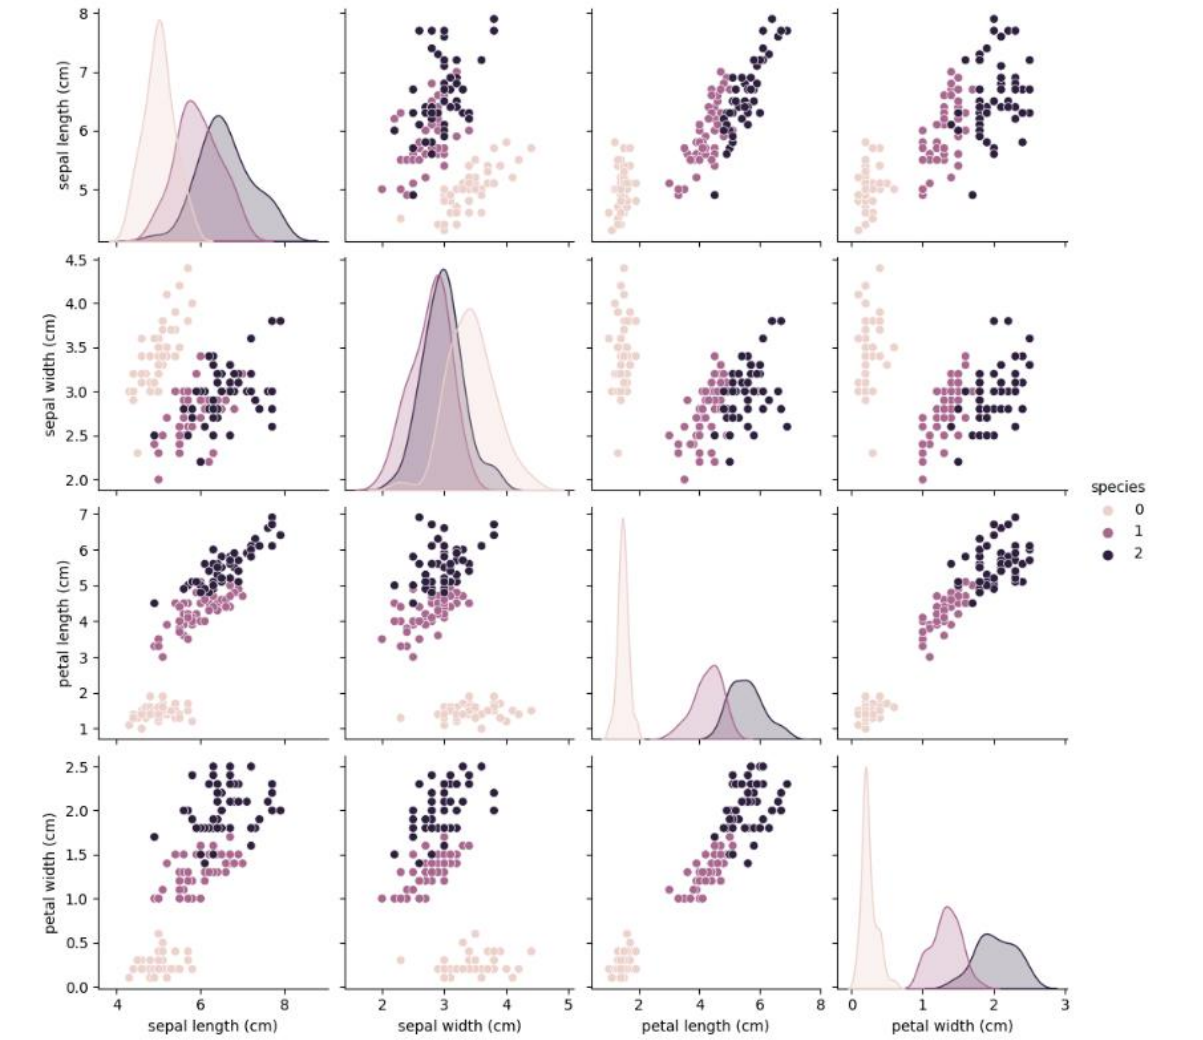
```



OUTPUT :

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

species	
0	0
1	0
2	0
3	0
4	0





Silhouette Score: 0.45994823920518635